

OM WANKAR

Pune, India · omgajananwankar123@gmail.com · +91 9325156044 · github.com/omwankar · linkedin.com/in/om-wankar

SUMMARY

Engineer building AI that stays tied to source context and only ships when the output holds up. Work spans LLM retrieval and confidence gates, evaluation and failure analysis, and the APIs and interfaces that turn model output into something a team can use. Python and TypeScript. Strongest fit: machine learning — testing approaches, measuring failures, and improving output quality.

EDUCATION

B.E. Electronics & Telecommunication — PVG's COET, Pune 2022 – 2026

EXPERIENCE

SDE-1 · Clarusto Technologies, Pune April 2025 – Present

- Built an ML credit-card reward tracker with Scikit-learn and feature engineering; measured 95% insight accuracy and identified \$50K in annual savings across 10K+ transactions.
- Designed an NLP transaction-categorization pipeline that raised engagement 25% and retention 40% through data-driven personalization.
- Built a real-time recommendation engine in pandas and NumPy that surfaces contextual reward suggestions and reduced user effort by 30%.

PROJECTS

Harbour Transcript Judge — evaluation without ground truth 2026

Python · failure taxonomy · calibration · github.com/omwankar/op04-harbour-transcript-judge

- Built a deterministic judge that reads an agent trajectory (customer message, tool calls, tool results) and returns a verdict, a confidence, and a failure category. The submitted judge makes no model call.
- Fit pass rates on 480 labelled development traces with case-grouped smoothing: balanced accuracy 0.867, MCC 0.801, AUROC 0.911, ECE 0.045, against a 0.50 always-pass baseline.
- Wrote the failure analysis: the 42 development errors sit in mixed “escalated instead of acting” leaves, only 3 of 480 errors occur at confidence ≥ 0.85 , and which rules transfer to unseen cases.

AI PR Review Agent — retrieval-grounded multi-agent review 2026

Python · FastAPI · LangGraph · hybrid retrieval · Next.js · github.com/omwankar/PR-reaview_system

- Built a pull-request review pipeline: verify the webhook, dedupe, enqueue, then four specialists (security, quality, tests, docs) that run only with retrieved code context.
- Hybrid retrieval fuses vector similarity and full-text search with reciprocal rank fusion. Findings are schema-valid; auto-post happens only when confidence is high and nothing is critical, otherwise a human queue.
- Added a golden-case regression gate, an append-only event log, and a budget check before model spend.

Scout — research agent that writes a sourced brief 2026

TypeScript · Claude · Playwright · Fastify · Next.js · scout-production-05ad.up.railway.app · github.com/omwankar/scout

- Shipped a browser agent that takes a plain-English research goal, searches and browses public pages, and finishes as a validated competitive brief with sources, exportable as Markdown or JSON.
- Streams thoughts, actions, and screenshots live. A person can stop the run, approve navigations, or steer it in plain language mid-flight. Deployed on Railway.

SKILLS

Languages — Python, TypeScript, JavaScript, SQL

AI / ML — LLM APIs (Claude, OpenAI), retrieval and embeddings, LangGraph, evaluation (accuracy, MCC, AUROC, calibration), Scikit-learn, TF-IDF, pandas, NumPy

Product — FastAPI, Next.js, React, Node.js, Playwright, Supabase / Postgres, REST APIs

Tools — Git, Docker, GitHub Actions, Railway, Vercel